

CS 61C Fall 2021

Part B: More Instructions

In this part, you will expand your CPU to support more instructions and pipelining.

You can implement the instructions in any order you want (see the appendix for a full list of instructions and subcircuit definitions). If you'd like a bit more guidance, this spec helps you build up your CPU with small groups of instructions at a time.

Task 6: I-type Instructions

The instructions you need to implement for this task are listed below:

Instruction	Type	Opcode	Funct3	Funct7/Immediate	Operation
addi rd, rs1, imm	I	0x13	0x0		$R[rd] \leftarrow R[rs1] + \text{imm}$
slli rd, rs1, imm			0x1	0x00	$R[rd] \leftarrow R[rs1] \ll \text{imm}$
slti rd, rs1, imm			0x2		$R[rd] \leftarrow (R[rs1] < \text{imm}) ? 1 : 0$
xori rd, rs1, imm			0x4		$R[rd] \leftarrow R[rs1] \wedge \text{imm}$
srli rd, rs1, imm			0x5	0x00	$R[rd] \leftarrow R[rs1] \gg \text{imm}$

srai rd, rs1, imm			0x5	0x20	$R[rd] \leftarrow R[rs1] \gg \text{imm}$
ori rd, rs1, imm			0x6		$R[rd] \leftarrow R[rs1] \mid \text{imm}$
andi rd, rs1, imm			0x7		$R[rd] \leftarrow R[rs1] \& \text{imm}$

Recall that you already implemented `addi` in Part A. These I-type instructions use the same datapath as `addi`, so your datapath already supports them!

Control Logic

As you add logic to support more instructions in the next few tasks, you will need to add control logic to enable the relevant datapath components depending on the instruction being executed.

The control logic should take the instruction bits and output all the control signals needed to execute that instruction. There are two main approaches to implementing this logic:

1. **Suggested: Read-only memory (ROM):** Hard-code a table of control signals. The relevant bits of the instruction are used to index into the table and find the control bits. See "Control Logic: Getting Started with ROM" for more information.
Pros: Fewer wires, more hard-coding (makes the control logic easier to debug).
Cons: ROM is common for CISC ISAs (like x86) but is less common for RISC ISAs (like RISC-V). Newer to this project (and may have some bugs we haven't discovered yet).
2. **Alternate: Hard-wired control:** Use logic gates (e.g. AND/OR/NOT gates, MUXes) to calculate the control bits from the instruction.
Pros: Used in real-world RISC-V systems, used in all past semesters.
Cons: More wires (harder to debug). Creating boolean equations with logic gates to calculate control bits generally takes longer than

hardcoding bits in the ROM.

Regardless of which approach you choose, you should modify `control_logic.circ` in each task to implement your control logic.

Here is a summary of the control signals you should implement:

Signal	Bit Width	Purpose
<code>PCSel</code>	1	Selects the ALU input for all SB-type instructions where the branch is taken (according to the branch comparator output). Selects the PC+4 input for all other instructions.
<code>ImmSel</code>	1	Selects the instruction format so the immediate generator can extract the immediate correctly.
<code>RegWEn</code>	1	1 if the instruction writes to a register, and 0 otherwise.
<code>BrUn</code>	1	1 if the branch instruction is unsigned, and 0 if the branch instruction is signed. Don't care for all other instructions.
<code>ASel</code>	1	Selects whether to send the data in <code>rs1</code> or the PC to the ALU.
<code>BSe1</code>	1	Selects whether to send the data in <code>rs2</code> or the immediate to the ALU.
<code>ALUSel</code>	4	Selects the correct operation for the ALU.
<code>MemRW</code>	1	1 if the instruction writes to memory, and 0 otherwise.
<code>WBSe1</code>	2	Selects whether to write the memory read from DMEM, the ALU output, or PC+4 to <code>rd</code> .

Control Logic: Getting Started with ROM

- `control-logic.circ` already has an empty ROM. To complete the control logic using ROM, you need to do the following:

▼ Details

We've created [a spreadsheet](#) to help you fill in the ROM. To generate data for the ROM, you should do the following: 1. Make a copy of the spreadsheet 2. For each instruction, fill in all columns under the "Control Signals" header (RegWEn through CSRWen). Enter binary digits without the leading `0b` (Ex: type `01` instead of `0b01`). If the upper-right corner of the control signal's cell turns red, or if you see a warning when you hover over the cell, your cell contents are invalid. Common sources of errors include entering a control signal that has too many or too few bits, or entering characters other than `0` or `1` in a control logic cell.

▼ Details

After you've filled out the spreadsheet from step 1, look for the column named "Rom Output". Copy the data in this column (not including the headers). Then, navigate to `control-logic.circ` and click on the ROM. In the properties tab on the left-hand sidebar, click on "(click to edit)" next to the "Contents" label. Click on the upper-left-most data cell (which will be a collection of 4 hex digits without a preceding `0x`) and paste your values from earlier into the ROM; you should see the ROM control bits change to the values from your spreadsheet. Click "Close Window" to exit the ROM programming view.

▼ Details

To select a specific set of control signals for an instruction, a unique input for each instruction must be passed into the ROM. The "ROM Input" header on the spreadsheet indicates the input that must be passed into the ROM for each instruction. For example, your control logic, when passed in an `addi` instruction, needs to pass `15` (more precisely `0b1111`) into the ROM so the correct entry of control signals can be selected.

To map instructions to their specific input, a Priority Encoder might be helpful! See the Example under "Tips: hard-wired control" for more information.

▼ Details

The `PCSel` control signal cannot be encoded in the ROM since it depends on runtime calculations performed by the processor. To

complete the control logic, write hard-wired control to drive the PCSel output.

- We do not suggest modifying the inputs or outputs to the existing starter `control_logic` circuit if you are using a ROM-based implementation, since the ROM and spreadsheet cannot be easily modified to output additional control signals. If you would like to add additional control signals, we suggest using hard-wired control logic to calculate them, like you did for PCSel.

Task 6.1: Control Logic

Refer to the control logic section above for instructions.

Testing and Debugging

Before you begin, keep in mind that debugging and testing should be conducted as you are implementing your CPU. It is **A LOT** easier to debug one instruction type at a time versus doing them all together. If you prefer a video format, take a look at our videos on [creating tests](#) and [debugging](#). Please note that the debugging video uses an older version of the test generation script; refer to it only for debugging and use the "creating tests" videos as a guide on how to create tests.

Here are the same instructions in written form:

We've included a script (`tools/create-test.py`) that uses Venus to help you generate test circuits from RISC-V assembly! The process for writing custom tests is as follows:

1. Come up with a test, and write the RISC-V assembly instructions for that test, saving them in a file ending in `.s` in the `tests/part-b/custom/inputs/` folder. The name of this file will be the name of your test. Repeat if you have more tests.
 - e.g. `tests/part-b/custom/inputs/sll-slli.s` and `tests/part-b/custom/inputs/beq.s`

Note: IMEM and DMEM are separate in Logisim, but combined in Venus. This means that if you write assembly code that tries to access memory overlapping with instructions, Venus will throw an error. Since counting exactly how many instructions your assembly code requires, and multiplying that by 4 can be annoying, we suggest you load/store using addresses greater

than 0x3E8 (leaving space for 1000 bytes/250 instructions), and increase this offset if you have more instructions.

2. Generate the test circuits for your tests using `create-test.py`:

```
$ python3 tools/create-test.py tests/part-b/custom/inputs/sll-
```

Reminder: if you want to regenerate *everything*, you can take advantage of Bash globs:

```
$ python3 tools/create-test.py tests/part-b/custom/inputs/*.s
```

This should generate a couple new files to go with your assembly file:

```
tests/part-b/custom/:
- cpu-<TEST_NAME>.circ                # The
- inputs/<TEST_NAME>.s                # The
- reference-outputs/cpu-<TEST_NAME>-pipelined-ref.out # The
- reference-outputs/cpu-<TEST_NAME>-ref.out          # The
```

3. Now you can run the tests you just wrote!

```
$ python3 test.py tests/part-b/custom/sll-slli.circ tests/part-
```

Reminder: you can run all tests in a directory:

```
$ python3 test.py tests/part-b/custom/
```

By default, the number of cycles for a test will be just enough for all instructions in the test, as well as extra cycles for the register writeback and pipelining. If you wish to override this and simulate your code for a certain number of cycles, you can use the `--cycles` flag:

```
$ python3 tools/create-test.py --cycles 10 tests/part-b/custom
```

4. Each test will print out either:

```
PASSED test: <TEST_NAME> # Your student-output/<TEST_NAME>-stu
FAILED test: <TEST_NAME> Your student-output/<TEST_NAME>-stude
```

5. If any of your tests fail, you will need to debug it more closely in Logisim

To debug any tests or see how the circuit is working for each instruction:

1. First find out where the output of your CPU implementation diverges from the staff solution. You can compare them using the diff command against the reference output and your output.

```
$ python3 tools/format-output.py tests/part-b/custom/reference
$ python3 tools/format-output.py tests/part-b/custom/student-c
$ diff reference.out student.out
```

2. Open the generated .circ file using Logisim. You will now see that test harness and instruction ROM. You can now use that information from your diff to step through the circuit to where you know your circuit is malfunctioning.
3. Trace your wires and check outputs that are malfunctioning. Consider all possibilities:

- If regA or regB is coming out wrong, maybe it was never set correctly
- If DMEM read is coming out wrong, maybe it was never set correctly
- If ALU output is wrong, maybe the control logic is malfunctioning
- If the instruction binary is wrong, maybe branching was not implemented correctly
- etc. etc., an important part of this project is to think about picking holes in your design and solving problems.

Task 6.2: Testing and Debugging

Refer to the testing and debugging section above for instructions.

Task 7: R-type Instructions

The instructions you need to implement for this task are listed below:

Instruction	Type	Opcod	Funct3	Funct7/Immediate	Operation
add rd, rs1, rs2	R	0x33	0x0	0x00	$R[rd] \leftarrow R[rs1] + R[rs2]$

mul rd, rs1, rs2		0x0	0x01	$R[rd] \leftarrow (R[rs1] * R[rs2])[31:0]$
sub rd, rs1, rs2		0x0	0x20	$R[rd] \leftarrow R[rs1] - R[rs2]$
sll rd, rs1, rs2		0x1	0x00	$R[rd] \leftarrow R[rs1] << R[rs2]$
mulh rd, rs1, rs2		0x1	0x01	$R[rd] \leftarrow (R[rs1] * R[rs2])[63:32]$
mulhu rd, rs1, rs2		0x3	0x01	(unsigned) $R[rd] \leftarrow (R[rs1] * R[rs2])[63:32]$
slt rd, rs1, rs2		0x2	0x00	$R[rd] \leftarrow (R[rs1] < R[rs2]) ? 1 : 0$ (signed)
xor rd, rs1, rs2		0x4	0x00	$R[rd] \leftarrow R[rs1] \wedge R[rs2]$
srl rd, rs1, rs2		0x5	0x00	(unsigned) $R[rd] \leftarrow R[rs1] >> R[rs2]$
sra rd, rs1, rs2		0x5	0x20	(signed) $R[rd] \leftarrow$

					$R[rs1] > > R[rs2]$
or rd, rs1, rs2			0x6	0x00	$R[rd] \leftarrow R[rs1] \mid R[rs2]$
and rd, rs1, rs2			0x7	0x00	$R[rd] \leftarrow R[rs1] \& R[rs2]$

Task 7.1: Datapath

Modify your datapath in `cpu.circ` so that it can support R-type instructions.

If you're stuck, read further for some guiding questions. As with Task 4, it may help to think about each of the five stages for executing an instruction.

▼ Instruction Fetch: How do R-type instructions affect the program counter?

R-type instructions always increment the program counter by 4 to fetch the next instruction, just like the `addi` instruction from Part A. This means we don't need to modify the program counter implementation for this task.

▼ Instruction Decode: What do we need to read from the register file?

R-type instructions have two source registers, `rs1` and `rs2`, that we need to read from the register file. In Part A, you split the bits from the instruction corresponding to the `rs1` address and passed them to the regfile. Now, you should also split the bits from the instruction corresponding to the `rs2` address and pass them to the regfile.

▼ Execute: What two data values (`A` and `B`) should an R-type instruction input to the ALU?

R-type instructions pass the register values from the regfile into the ALU. In Part A, you already passed the first register value $R[rs1]$ into the first input of the ALU. However, for the `addi` instruction, the second input of the ALU is an immediate. Since you want to support

both R-type instructions and the `addi` instruction, you should use a multiplexer to select which input will be inputted to the ALU.

The select bit of this multiplexer is `BSe1`. You will implement the logic for determining `BSe1` from the instruction bits in the control logic later in this task.

▼ Memory: Do R-type instructions write to memory?

R-type instructions do not write to memory (they write to a register on the CPU, which is different from memory). This means we don't need to modify DMEM for this task.

▼ Write back: What data is the R-type instruction writing, and where is the instruction writing this data to?

R-type instructions take the result of the computation (from the ALU output) and write the result to the register `rd`. In Part A, you already implemented logic to write the ALU output into a destination register.

Task 7.2: Control Logic

Refer to the control logic section above for instructions.

Task 7.3: Testing and Debugging

Refer to the testing and debugging section above for instructions.

Task 8: SB-type Instructions

The instructions you need to implement for this task are listed below:

Instruction	Type	Opcode	Funct3	Funct7/Immediate	Operation
beq rs1, rs2, offset	SB	0x63	0x0		if($R[rs1] == R[rs2]$) $PC \leftarrow PC + \{offset, 1b0\}$
bne rs1, rs2, offset			0x1		if($R[rs1] != R[rs2]$)

					PC \leftarrow PC + {offset, 1b0}
blt rs1, rs2, offset			0x4		if(R[rs1] < R[rs2] (signed)) PC \leftarrow PC + {offset, 1b0}
bge rs1, rs2, offset			0x5		if(R[rs1] >= R[rs2] (signed)) PC \leftarrow PC + {offset, 1b0}
bltu rs1, rs2, offset			0x6		if(R[rs1] < R[rs2] (unsigned)) PC \leftarrow PC + {offset, 1b0}
bgeu rs1, rs2, offset			0x7		if(R[rs1] >= R[rs2] (unsigned)) PC \leftarrow PC + {offset, 1b0}

Task 8.1: Branch Comparator

Fill in the branch comparator subcircuit in `branch-comp.circ`. This subcircuit takes two inputs and outputs the result of comparing the two inputs. We will use the output later for implementing branches.

Signal Name	Direction	Bit Width	Description
rs1	Input	32	First value to compare
rs2	Input	32	Second value to compare
BrUn	Input	1	1 when an unsigned comparison is wanted, and 0 when a signed comparison is wanted
BrEq	Output	1	Set to 1 if the two values are equal
BrLt	Output	1	Set to 1 if the value in rs1 is less than the value in rs2

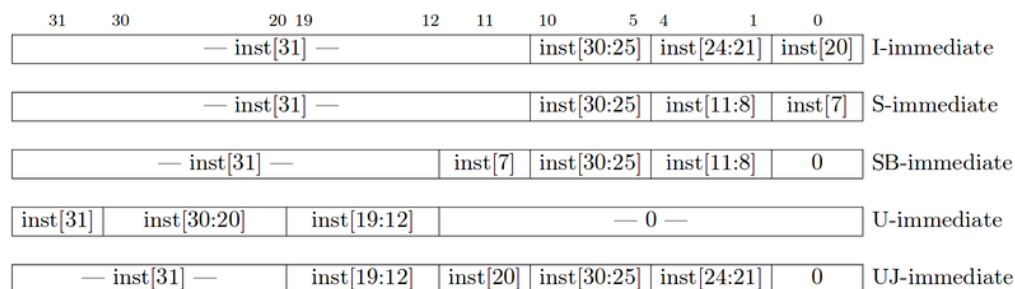
Task 8.2: Immediate Generator

Edit the immediate generator in `imm-gen.circ` so that it can generate immediates for SB-type instructions in addition to immediates for I-type instructions (which you implemented in Part A).

Recall that the bits of the immediate are stored in different bits of the instruction, depending on the type of the instruction. The `ImmSel` signal, which you will implement in the control logic, will determine which type of immediate this subcircuit should generate.

Also, recall that all immediates are 32 bits and sign-extended.

The immediate storage formats are listed below:



Task 8.3: Datapath

Modify your datapath in `cpu.circ` so that it can support SB-type instructions.

If you're stuck, read further for some guiding questions. As with Task 4, it may help to think about each of the five stages for executing an instruction.

▼ Instruction Fetch: How do SB-type instructions affect the program counter?

Recall that branching instructions add an immediate to the current value of PC. If the branch is taken, the PC changes to be the result of this addition. If the branch is not taken, or the instruction is not an SB-type instruction, then PC changes to PC+4 (just like in the previous tasks). We will implement this in the write-back stage.

▼ Instruction Decode: What do we need to read from the register file?

SB-type instructions have two source registers, `rs1` and `rs2`, that we need to read from the register file. In the previous task, you already implemented reading `rs1` and `rs2` for R-type instructions.

▼ Execute: What two data values (`A` and `B`) should an SB-type instruction input to the ALU?

SB-type instructions use the ALU to add an immediate to PC. You will need to add a multiplexer so that the ALU can receive either PC or the value in `rs1`, depending on the instruction being executed. The select bit of this multiplexer is `ASel`. In the previous tasks, you already implemented sending an immediate to the ALU.

▼ Memory: Do SB-type instructions write to memory?

SB-type instructions do not write to memory. This means we don't need to modify DMEM for this task.

▼ Write back: What data is the SB-type instruction writing, and where is the instruction writing this data to?

SB-type instructions take the result of the addition (PC + immediate, from the ALU output) and might write the result to PC (depending on if the branch is taken). You should use a multiplexer to select which value will be written to PC.

The select bit of this multiplexer is `PCSel`. You will implement the logic for determining `PCSel` from the instruction bits in the control logic.

Task 8.4: Control Logic

Refer to the control logic section above for instructions.

Task 8.5: Testing and Debugging

Refer to the testing and debugging section above for instructions.

Task 9: Loading and Storing

The instructions you need to implement for this task are listed below:

Instruction	Type	Opcode	Funct3	Funct7/Immediate	Operation
lb rd, offset(rs1)	I	0x03	0x0		$R[rd] \leftarrow \text{SignExt}(\text{Mem}(R[rs1] + \text{offset, byte}))$
lh rd, offset(rs1)			0x1		$R[rd] \leftarrow \text{SignExt}(\text{Mem}(R[rs1] + \text{offset, half}))$
lw rd, offset(rs1)			0x2		$R[rd] \leftarrow \text{Mem}(R[rs1] + \text{offset, word})$
sb rs2, offset(rs1)	S	0x23	0x0		$\text{Mem}(R[rs1] + \text{offset}) \leftarrow R[rs2][7:0]$
sh rs2, offset(rs1)			0x1		$\text{Mem}(R[rs1] + \text{offset}) \leftarrow R[rs2][15:0]$
sw rs2, offset(rs1)			0x2		$\text{Mem}(R[rs1] + \text{offset}) \leftarrow R[rs2]$

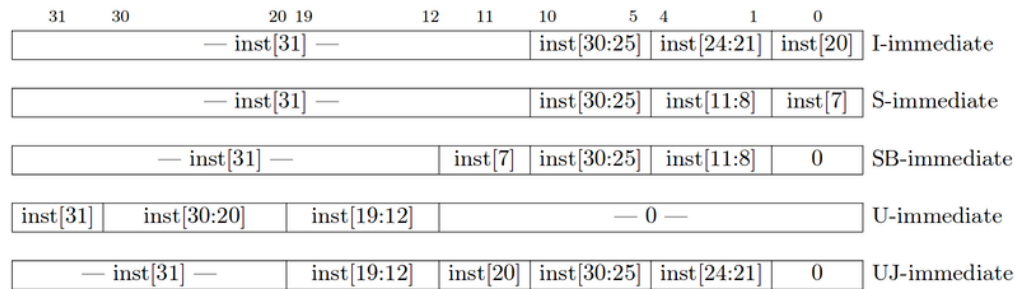
Task 9.1: Immediate Generator

Edit the immediate generator in `imm-gen.circ` so that it can generate immediates for S-type instructions in addition to all the instruction types from previous tasks.

Recall that the bits of the immediate are stored in different bits of the instruction, depending on the type of the instruction. The **ImmSel** signal, which you will implement in the control logic, will determine which type of immediate this subcircuit should generate.

Also, recall that all immediates are 32 bits and sign-extended.

The immediate storage formats are listed below:



Conceptual Overview: Memory Unit

Since these instructions load from memory and store to memory, you will need to interact with the memory subcircuit (already implemented for you).

Signal Name	Direction	Bit Width	Description
WriteAddr	Input	32	Address to read/write to in Memory
WriteData	Input	32	Value to be written to Memory
Write_En	Input	4	The write mask for instructions that write to Memory and zero otherwise
CLK	Input	1	Driven by the clock input to the CPU
ReadData	Output	32	Value of the data stored at the specified address

Due to Logisim size limitations, the memory unit only uses the lower 16 bits of the provided address, discarding the upper 16 bits. This means that the memory can only store 2^{16} bytes of data. The provided tests

will always set the upper 16 bits of addresses to 0, and your tests should avoid using the upper 16 bits when interacting with memory.

The memory unit will always zero out the bottom two bits of the provided address and read 4 bytes starting at this modified address. In other words, the provided address will be rounded down to the nearest multiple of 4, and 4 bytes will be read in that range of 4 bytes.

For example, if the input address `WriteAddr` is `0x00001007`, the bottom two bits will be zeroed out to make `0x00001004`, and the output `ReadData` will be the 4 bytes at addresses `0x00001004`, `0x00001005`, `0x00001006`, and `0x00001007`.

The same process of zeroing out the bottom two bits applies for writing data as well, but the 4-bit `Write_En` input lets you select which of the 4 bytes in the `WriteData` input are written to the 4 bytes at the zeroed-out `WriteAddr` address. Each bit of this write mask enables writing to the corresponding byte of the word.

For example, if the input address `WriteAddr` is `0x00001007`, the write mask `Write_En` is `0b0010`, and the `WriteData` input is `0x11223344`, then the byte `0x33` will be written to `0x00001006`. If `Write_En` was changed to `0b0011`, then the byte `0x44` would also be written to `0x00001007`.

Conceptual Overview: Alignment

In this project, all memory accesses will be *aligned*. This means that a single load or store instruction will never cross a word boundary in memory.

All `lw` and `sw` instructions will use memory addresses that end in `0b00` (accessing `0b00`, `0b01`, `0b02`, and `0b03` in memory).

All `lh` and `sh` instructions will use memory addresses that end in either `0b00` (accessing `0b00` and `0b01`) or `0b10` (accessing `0b10` and `0b11`).

You should not implement any unaligned memory accesses in this project.

Task 9.2: Datapath

Modify your datapath in `cpu.circ` so that it can support loads and stores.

You should provide an address input `WriteAddr` to DMEM. Remember that the ALU calculates this address by adding the address in `rs1` and the offset immediate.

For load instructions, you should also add functionality in the write-back stage so that the data outputted by DMEM (and processed by your logic in Task 9.3) is written back to the `rd` register.

Task 9.3: Load

Remember that the address input `WriteAddr` will have its bottom two bits zeroed out. The DMEM will then output `ReadData`, which contains 4 bytes starting at this modified address. Implement logic in `cpu.circ` to extract and, if needed, sign-extend the relevant byte(s) from `ReadData`. Only the relevant byte(s) with appropriate sign-extending should be written back to the `rd` register.

For completeness, a table of scenarios you need to handle is provided below:

Instruction	Type	Opcode	Funct3	Bottom 2 bits of address	Value to put in rd
lb rd, offset(rs1)	I	0x03	0x0	0b00	SignExt(ReadData[7:0])
				0b01	SignExt(ReadData[15:8])
				0b10	SignExt(ReadData[23:16])
				0b11	SignExt(ReadData[31:24])
lh rd, offset(rs1)			0x1	0b00	SignExt(ReadData[15:0])
				0b10	SignExt(ReadData[31:16])
lw rd, offset(rs1)			0x2	0b00	ReadData

Task 9.3: Store

Implement logic in `cpu.circ` to set up the `WriteData` input and `Write_En` write mask to correctly store the data in `rs2` to memory.

For completeness, a table of scenarios you need to handle is provided below:

Instruction	Type	Opcode	Funct3	Bottom 2 bits of address	WriteData	Write_En
sb rs2, offset(rs1)	S	0x23	0x0	0b00	{24'b0, R[rs2] [7:0]}	0001
				0b01	{16'b0, R[rs2] [7:0], 8'b0}	0010
				0b10	{8'b0, R[rs2] [7:0], 16'b0}	0100
				0b11	{R[rs2] [7:0], 24'b0}	1000
sh rs2, offset(rs1)			0x1	0b00	{16'b0, R[rs2] [15:0]}	0011
				0b10	{R[rs2] [15:0], 16'b0}	1100
sw rs2, offset(rs1)			0x2	0b00	R[rs2]	1111

Note that for any non-store instruction (i.e. when your `MemRW` control signal is 0), the `Write_En` write mask should be set to `0000`.

Task 9.4: Control Logic

Refer to the control logic section above for instructions.

Task 9.5: Testing and Debugging

Refer to the testing and debugging section above for instructions.

Task 10: All Other Instructions

The instructions you need to implement for this task are listed below:

Instruction	Type	Opcode	Funct3	Funct7/Immediate	Operation
auipc rd, offset	U	0x17			$R[rd] \leftarrow PC + \{\text{offset}, 12b0\}$
lui rd, offset		0x37			$R[rd] \leftarrow \{\text{offset}, 12b0\}$
jal rd, imm	UJ	0x6f			$R[rd] \leftarrow PC + 4$ $PC \leftarrow PC + \{\text{imm}, 1b0\}$
jalr rd, rs1, imm	I	0x67	0x0		$R[rd] \leftarrow PC + 4$ $PC \leftarrow R[rs1] + \{\text{imm}\}$

Task 10.1: Datapath

Modify your datapath in `cpu.circ` so that it can support these instructions. Most of these instructions are already supported by your datapath so far.

To support `jalr`, you should connect `PC+4` to your multiplexer in the write-back stage so that `PC+4` can be written back to `rd`.

Task 10.2: Control Logic

Refer to the control logic section above for instructions.

Task 10.3: Testing and Debugging

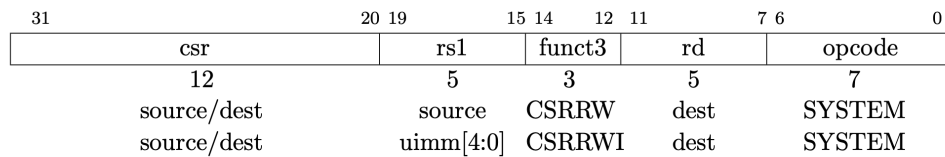
Refer to the testing and debugging section above for instructions.

Task 11: CSRW

The instructions you need to implement for this task are listed below:

Instruction	Type	Opcode	Funct3	Funct7/Immediate	Operation
csw rd, csr, rs1	I	0x73	0x1		CSR[csr] ← R[rs1]
cswi rd, csr, uimm			0x5		CSR[csr] ← {uimm}

The instruction format for these instructions is shown below:



Conceptual Overview: Control Status Registers (CSRs)

Control Status Registers (CSRs) are used to hold additional information about the instructions being executed, such as debugging information or testing flags. These are unrelated to the registers and memory from the previous tasks.

To fully test your implementation, the project autograder requires one CSR (`tohost = 0x51E`). This means that for this project, the `csr` field will always be `0x51E` in this project.

Also, for this project, `rd` will always be 0, so you don't need to worry about CSR instructions writing to a register.

Task 11.1: Immediate Generator

Edit the immediate generator in `imm-gen.circ` so that it can generate immediates for the `csrwi` instruction (in addition to all the previous immediates you've implemented).

Note that the immediate is in bits 19-15 (labeled `uimm[4:0]`) and should be zero-extended, not sign-extended.

Task 11.2: Datapath

Modify your datapath in `cpu.circ` so that it can support these instructions.

You will need to interact with the CSR subcircuit (already implemented for you).

Signal Name	Direction	Bit Width	Description
CSR_address	Input	12	Input CSR register address
CSR_din	Input	32	Value to write into specified CSR register
CSR_WEn	Input	1	Write enable (from control logic)
clk	Input	1	Clock input
tohost	Output	32	Output of the <code>tohost</code> register

You should connect `CSR_address` to the 12-bit `csr` field in the instruction.

For `csrwi` instructions, `CSR_din` should be the data in `rs1`. For `csrwi` instructions, `CSR_din` should be the immediate from the immediate generator.

Task 11.3: Control Logic

Refer to the control logic section above for instructions.

Note that there are two additional control signals for the CSR instructions:

Signal	Bit Width	Purpose
CSRSe1	1	Selects whether the instruction is <code>csrwr</code> or <code>csrwi</code> . Don't care for all other instructions.
CSRWen	1	Selects whether the instruction is a CSR instruction. 1 for <code>csrwr</code> and <code>csrwi</code> , 0 otherwise.

Task 12: Pipelining

In this task, you will implement a 2-stage pipeline in your CPU:

1. **Instruction Fetch:** An instruction is fetched from the instruction memory.
2. **Execute:** The instruction is decoded, executed, and committed (written back). This is a combination of the remaining four stages of a classic five-stage RISC-V pipeline (ID, EX, MEM and WB).

Some things to consider:

- Will the IF and EX stages have the same or different `PC` values?
- Do you need to store the `PC` between the pipelining stages?
- What hazards are present in this two-stage pipeline?

Note: During the first cycle, the instruction register sitting between the pipeline stages won't contain an instruction loaded from memory. What should the second stage do? Luckily, Logisim automatically sets registers to zero on reset, so the instruction pipeline register will automatically start with a `nop`! If you wish, you can depend on this behavior of Logisim.

Since your CPU will support branch and jump instructions, you'll need to handle control hazards that occur when branching.

- The instruction immediately after a branch or jump should not be executed if a branch is taken. This makes your task a bit more complex. By the time you have figured out that a branch or jump is in the execute stage, you have already accessed the instruction memory and pulled out (possibly) the wrong instruction. Therefore, you will need to flush the instruction that is being fetched (next instruction) if the instruction under execution (current instruction) is a taken branch or jump.

- Instruction flushing for this project must be accomplished by MUXing a `nop` into the instruction stream and sending the `nop` into the Execute stage instead of using the fetched instruction. You can use `addi x0, x0, 0 (0x00000013)`, for this purpose; other `nop` instructions will work too. You should flush only if a branch is taken (do not flush if it is not taken). You should always flush the next instruction when jumping.
- Note: you should not solve this issue by calculating branch offsets in the IF stage. We compare your output against a reference output that uses `nop`, so a solution that doesn't use `nop` may not match the reference output even if it is a conceptually correct solution.

Some more things to consider:

- To MUX a `nop` into the instruction stream, do you place it *before* or *after* the instruction register?
- What address should be requested next while the EX stage executes a `nop`? Is this different than normal?

Testing

We've provided some basic sanity tests for your pipelined CPU in the `tests/part-b/sanity/` directory (same tests as in Task 6). You can run these with:

```
$ python3 test.py --pipelined tests/part-b/sanity/
```

Note: since your CPU is pipelined at this point, you need to run the pipelined tests using the `--pipelined` (or `-p`) flag. If you run the single-cycle tests (i.e. omit the `--pipelined` flag) after pipelining your CPU, your CPU should now fail those tests! Think about why this happens...

Similarly, you can also run the pipelined version of your custom tests:

```
$ python3 test.py --pipelined tests/part-b/custom/
```

Note: because you're implementing a 2-stage pipelined processor and the first instruction writes on the rising edge of the second clock cycle, the effects of your instructions will have a 2 instruction delay. For example, let's look at the first instruction of `tests/part-b/sanity/inputs/addi.s`, `addi t0, x0, -1`. If you inspect the pipelined reference output (`tests/part-b/sanity/reference-output/cpu-addi-`

`pipelined-ref.out`), you'll see that `t0` doesn't show changes until the third cycle.

Refer to the section from Project 3A for more info on using these tests. Keep in mind that you're working with a pipelined circuit from this task onward.

Task 13: Part B README Update

Time to update your `README.md`! Once again, write down how you implemented your circuits and components for this part (including the various subcircuits you used), and explain the reasoning behind your design choices. In particular, we want to see:

- How you designed your control logic
- Advantages/Disadvantages of your design
- Best/Worst bug or design challenge you encountered, and your solution to it

Your additions to the README should be at least 512 characters (although something more than the bare minimum would be nice), but other than that feel free to get creative!

Submission and Grading

Submit your assignment to the Project 3B submission on Gradescope. Part B is worth 80% of your overall Project 3 grade.

- Unit Tests (35%)
- Test Coverage (10%)
- Integration (5%)
- Edge (29%)
- README (1%)

☐ Dark Mode